



Intro to Docker

Presented by **Travis Holton**

Administrivia

- Bathrooms
- Fire exits

This course

- Makes use of official Docker docs
- Based on latest Docker
- A mix of command line and theory
- Assumes no prior Docker knowledge
- Assumes familiarity with the linux shell
- Assumes we are using ubuntu 14.04 (trusty)

Aims

- Understand how to use Docker on the command line
- Understand how Docker works
- Learn how to integrate Docker with applications
- Learn ops and developers can use Docker to deploy applications
- Get people thinking about where they could use Docker

Setup

Fetch course resources

```
$ git clone \  
  https://github.com/catalyst-training/docker-introduction.git  
$ cd docker-introduction
```

```
$ cd ~/docker-introduction  
$ ls
```

- Slides for Reveal.js presentation
- docker-introduction.pdf
- Ansible setup playbook
- Sample code for some exercises

Ansible

- Some of the features we will be exploring require setup. We'll use ansible for that.
- Python based tool set
- Automate devops tasks
 - server/cluster management
 - installing packages
 - deploying code
 - managing config

Setup Ansible

```
$ git clone https://github.com/catalyst/catalystcloud-ansible.git
```

```
$ cd ~/catalystcloud-ansible  
$ ./install-ansible.sh  
.  
. <stuff happens>  
.  
$ source $CC_ANSIBLE_DIR/ansible-venv/bin/activate
```

- Installs python virtualenv with latest ansible libraries
- We'll be using this virtualenv for tasks throughout the course.

Setup Docker

- Follow instructions on website for installing
 - Docker Community Edition
 - docker-compose
- If you are using Ubuntu, use the ansible playbook included in course repo

```
$ cd docker-introduction
$ ansible-playbook -K ansible/docker-install.yml \
  -e ansible_python_interpreter=/usr/bin/python
```

- This playbook installs:
 - latest Docker *Community Edition*
 - docker - compose
 - Note: you might need to logout and login again

Fetch and run slides

```
$ docker run --name docker-intro -d --rm \
  -p 8000:8000 heytrav/docker-introduction-slides
```

Follow along with course slides: <http://localhost:8000>

Introduction to containers

What is containerization?

- A type of virtualization
- Difference from traditional VMs
 - Don't replicate entire OS, just bits needed for application
 - Run natively on host
- Key benefits:
 - More lightweight than VMs
 - Efficiency gains in storage, CPU
 - Portability

Lightweight

Virtualization

App A	App B
Bins / Libs	Bins / Libs
Guest OS	Guest OS
Hypervisor	
Host OS	
Server	

Docker

App A	App B
Bins / Libs	Bins / Libs
Docker engine	
Host OS	
Server	

Benefits of Containers: Resources

- Containers share a kernel
- Use less CPU than VMs
- Less storage. Container image only contains:
 - executable
 - application dependencies

Benefits of Containers: Decoupling

- Application stack not coupled to host machine
- Scale and upgrade services independently
- Treat services like cattle instead of pets



Benefits of Containers: Workflows

- Easy to distribute
- Developers can wrap application with libs and dependencies as a single package
- Easy to move code from development environments to production in easy and replicable fashion

Introduction to Docker

The Docker Platform

What is Docker?

High level

An open-source platform for creating, running, and distributing software *containers* that bundle software applications with all of their dependencies.

Low level

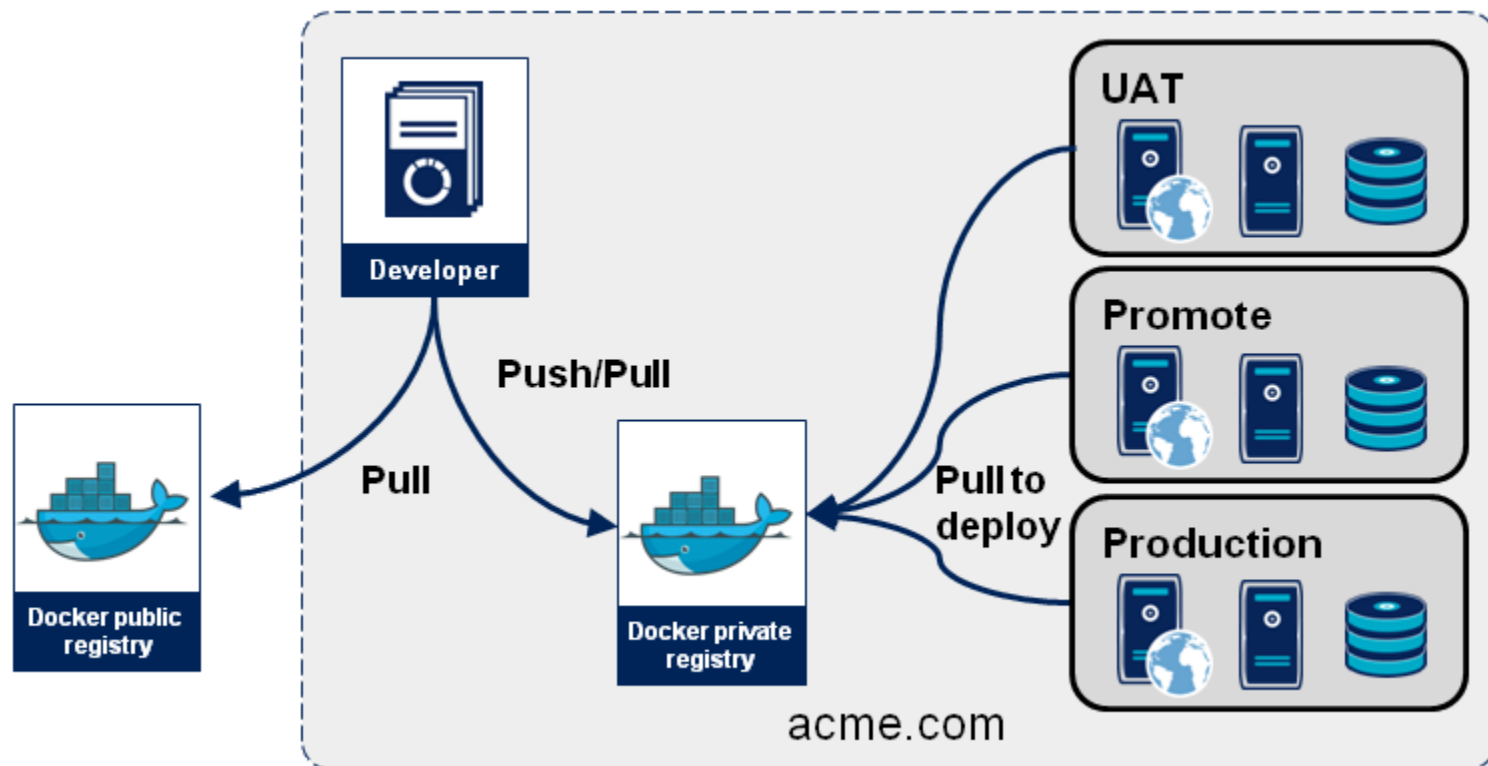
A command-line tool for programmatically defining the contents of a Linux container in code, which can then be versioned, reproduced, shared, and modified easily just as if it were the source code to a program

Docker popularity

- Linux containers are not new
 - FreeBSD Jails
 - LXC containers
 - Solaris Zones
- Docker is doing for containers what Vagrant did for virtual machines
 - Easy to create
 - Easy to distribute

Docker workflow

- Developer packages application and supporting components into image
- Developer/CI pushes image to private or public registry
- The image becomes the unit for distributing and testing your application.



Docker Portability

- Most modern operating systems
 - Linux (RHEL, CentOS, Ubuntu LTS, etc.)
 - OSX
 - Windows
- Lightweight Docker optimized Linux distributions (CoreOS, Project Atomic, RancherOS, etc.)
- Private clouds (OpenStack, Vmware)
- Public clouds (AWS, Azure, Rackspace, Google)

First Steps with Docker

Docker version

```
$ docker --version  
Docker version 17.12.0-ce, build c6d4123
```

Current version scheme similar to Ubuntu versioning:
YY.MM.#

Get command documentation

- Just typing **docker** returns list of commands
- Comprehensive online docs on **Docker website**

```
$ docker<ENTER>
```

```
Usage:  docker COMMAND
```

```
A self-sufficient runtime for containers
```

```
Options:
```

<code>--config string</code>	Location of client config files (default <code>~/.docker</code>)
<code>-D, --debug</code>	Enable debug mode
<code>--help</code>	Print usage

```
.
```

```
.
```


Basic client usage

docker **command** *[options]* *[args]*

- Calling any command with **--help** displays some docs

Exercise: View documentation for docker run

```
$ docker run --help
```

```
Usage:  docker run [OPTIONS] IMAGE [COMMAND] [ARG...]
```

Run a command **in** a **new** container

Options:

-- add -host list	Add a custom host- to -IP mapping (host:port:ip)
-a, --attach list	Attach to STDIN, STDOUT or STDERR
--blkio-weight uint16	Block IO (relative weight), between 1 and 1024
--blkio-weight-device list	Block IO weight (relative device weight)
.	
.	

Search for images

docker search [OPTIONS] TERM

Option	Argument	Description
-f, --filter	filter	Filter output based on conditions provided
--format	string	Pretty-print search using a Go template
--help		Print usage
--limit	int	Max number of search results (default 25)
--no-trunc		Don't truncate output

```
$ docker search hello-world
```



00:06



Pull an image from a registry

`docker pull [OPTIONS] NAME[:TAG]`

Option	Argument	Description
-a, --all-tags		Download all tagged images in the repository
--disable-content-trust		Skip image verification (default true)
--help		Print usage

```
$ docker pull hello-world
```



00:06



docker run

Run a command in a container from an image

docker run [options] *image* [*command*]

- docker run requires an *image* argument

Option	Argument	Description
-i		Keep STDIN open
-t		Allocate a tty
--rm		Automatically remove container on exit
-v	list	Mount a volume
-p	list	List of port mappings
-e, --env	list	Set environment variables
-d, --detach		Run container in background and print container ID
--link	list	Add link to another container
--name	string	Name for the container

These are just examples that we'll use in the course. See complete list with `docker run --help`

Run a simple container

```
$ docker run hello-world
```

```
> docker run hello-world  
Unable to find image 'hello-world:latest' locally
```

00:06

- The hello-world image was created by docker for instructional purposes. It just outputs a *hello world*-like message and exits.

Execute command in a container

`docker run image [command]`

```
$ docker run alpine ls
bin
dev
.
.
usr
var
$
```

- Docker starts container using alpine image
- The *alpine* image contains the **Alpine OS**, a very minimal Linux distribution.
- `[command]` argument is executed inside container
- Exits immediately
- A docker container only runs as long as it has a process (eg. a shell terminal or program) to run

Exercise: Start an *interactive* shell

docker run [options] alpine /bin/sh

Find [options] to make container run interactively

```
$ docker run -it alpine /bin/sh
```

```
→ ~ docker run -it alpi
```

- Docker starts alpine image
 - -i interactively
 - -t allocate a pseudo-TTY
- Runs shell command
- Execute commands inside container
- Exiting the shell stops the process and the container

List running containers

docker ps

```
$ docker ps
```

```
CONTAINER ID   IMAGE                                     ... NAMES
b3169acf49f8   alpine                                   ... adoring_edison
02aa3e50580c   heytrav/docker-introduction-slides     ... docker-intro
```

Note: by default docker will assign a random name to each container (i.e. *adoring_edison*).

Option	Argument	Description
-a, --all		Show all containers (default shows just running)
-f, --filter	filter	Filter output based on conditions provided
--format	string	Pretty-print containers using a Go template
--help		Print usage
--no-trunc		Don't truncate output

Exercise: Assign the name *myalpine* when running previous example container

Hint: `docker run -it <option> alpine`

```
$ docker run -it --name myalpine alpine /bin/sh
```

```
> docker ps
```

CONTAINER ID	IMAGE	...	NAMES
db1faf244e7a	alpine	...	myalpine
02aa3e50580c	heytrav/docker-introduction-slides	...	docker-intro

- Exit the shell
- Repeat using same name. What happens?

```
> docker run -it --name myalpine
```

Removing containers

```
docker rm name|containerID
```

Exercise: remove old myalpine container

```
$ docker rm myalpine
```

```
~  
> docker rm myalpine  
myalpine  
~  
>
```



00:06

If you pass the **--rm** flag to `docker run`, containers will be cleaned up when stopped.

Exercise: Run website in a container

```
$ docker run [OPTIONS] dockersamples/static-site
```

- Find values for [OPTIONS]:
 - Give it the name: *static-site*
 - Pass `AUTHOR="YOURNAME"` as environment variable
 - Map port 8081 to 80 internally (hint `8081:80`)
 - Cleans up container on exit

```
> docker run --name static-site --rm \
    -e AUTHOR="Trav" -p 8081:80 dockersamples/static-site
Unable to find image 'dockersamples/static-site:latest' locally
latest: Pulling from dockersamples/static-site
fdd5d7827f33: Pulling fs layer
a3ed95caeb02: Pulling fs layer
716f7a5f3082: Waiting
7b10f03a0309: Waiting
```

II 00:06

- Note: `docker run` implicitly pulls image if not available
- Try to exit using CTRL-C. What happens?

Stop a running container

```
docker stop name|containerID
```

Exercise: Stop the static-site container

- You actually have a couple options:
 - use the name you gave to the container

```
$ docker stop static-site
```

- use the CONTAINERID from `docker ps` output (will depend on your environment)

```
$ docker stop 25eff330a4e4
```

Exercise: Running a detached container

- Run static-site container like you did before, but add option to run in the background (i.e. *detached* state).

```
$ docker run --rm --name static-site -e AUTHOR="YOUR NAME" \  
-d -p 8081:80 dockersamples/static-site
```

```
> docker run --rm --name static-site -e AUTHOR="YOUR NAME" \  
-d -p 8081:80 dockersamples/static-site  
06ba2a841d43ad02a81e33f62561c87c5fb840aebcb0243e6b1a2c6d59a1e16d
```

```
> docker port s
```



00:06



View container logs

`docker logs [options] CONTAINER`

Option	Argument	Description
<code>--details</code>		Show extra details provided to logs
<code>-f, --follow</code>		Follow log output
<code>--help</code>		Print usage
<code>--since</code>	string	Show logs since timestamp (e.g. 2013-01-02T13:23:37) or relative (e.g. 42m for 42 minutes)
<code>--tail</code>	string	Number of lines to show from the end of the logs (default "all")
<code>-t, --timestamps</code>		Show timestamps

See [online documentation](#)

Exercise: view container logs for static-site container

```
> docker logs -f static-site
```



00:06



Note: Go to **localhost:8081** and refresh a few times

docker exec

docker **exec** [options] **CONTAINERID** [command]

- A way to interact with a running container
- Open a shell inside a running container.
- A bit like ssh'ing into a machine
- Can be useful for debugging
- See [online documentation](#)

Exercise: Check process list in static-site container

```
> docker exec -it static-s
```




00:06



List local images

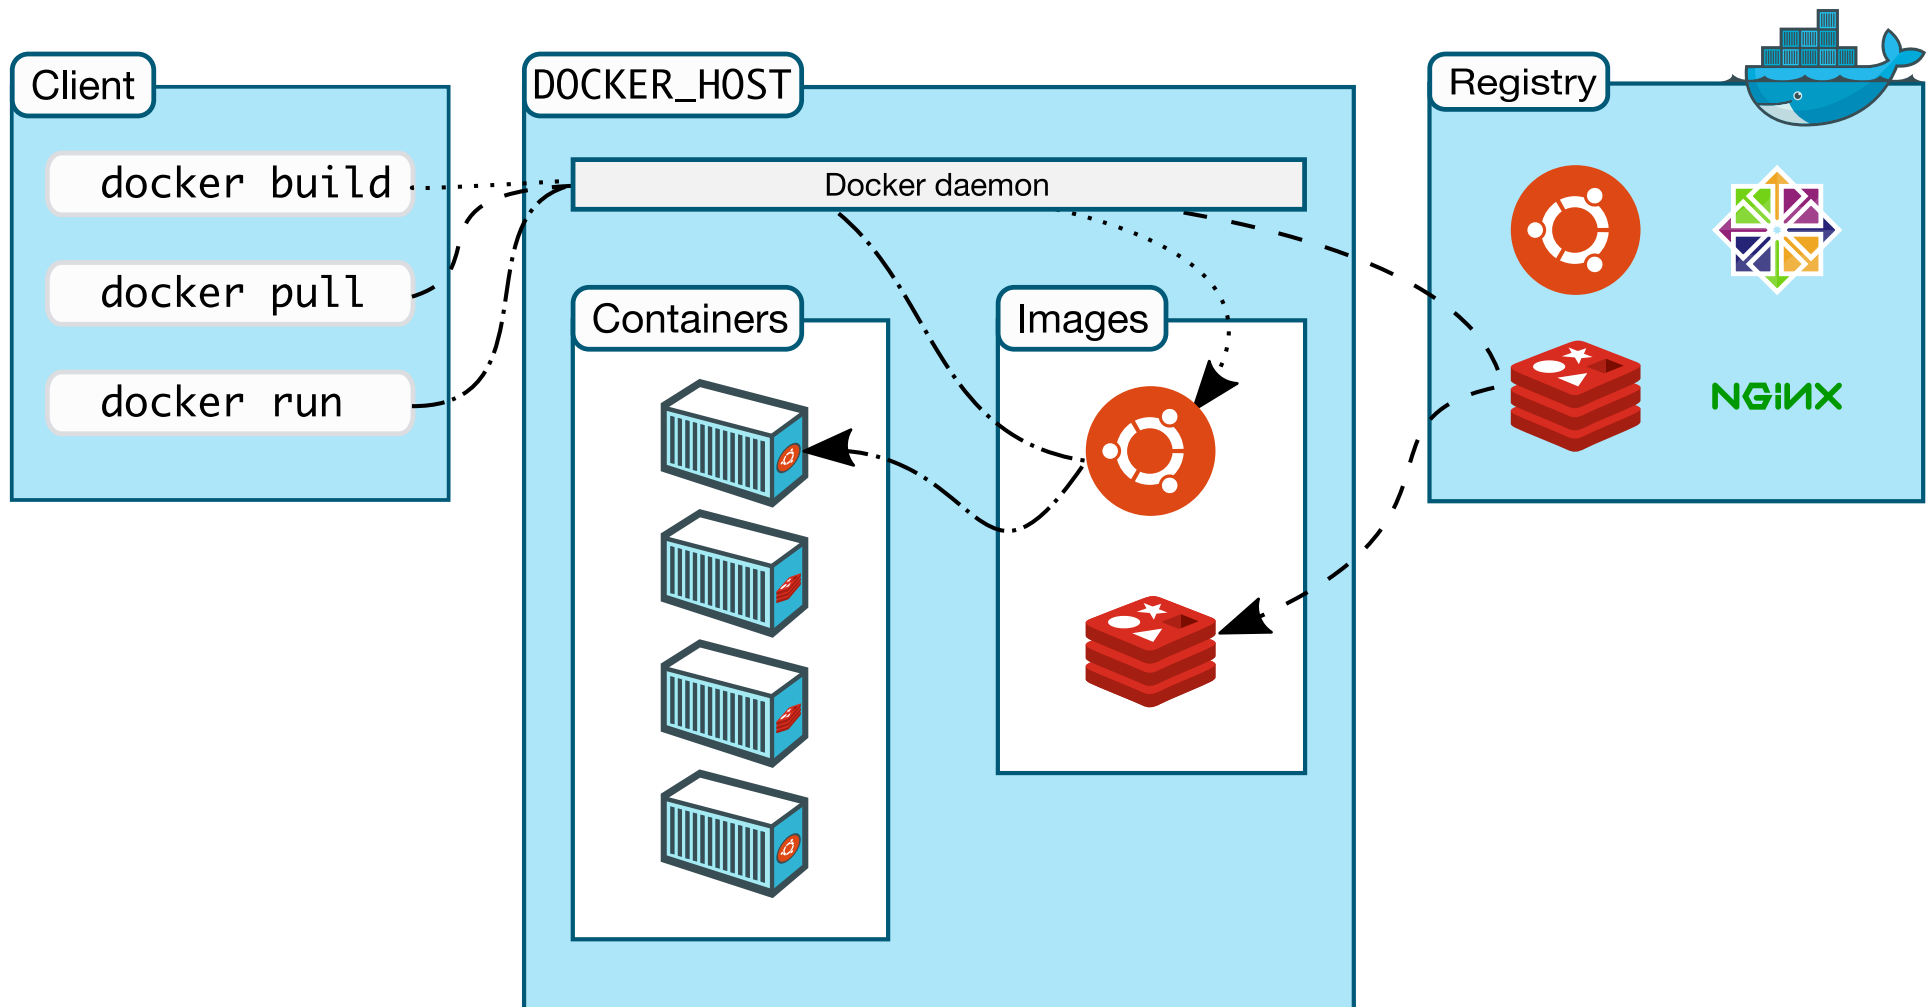
```
$ docker image ls
```



Behind the scenes

- User types docker commands
- Docker client contacts docker daemon
- Docker daemon checks if image exists
- Docker daemon downloads image from docker registry if it does not exist
- Docker daemon runs container using image

Docker architecture



How Docker works

Components of Docker

- Images
 - The build component
 - Distributable *artefact*
- Containers
 - The run component
- Registries
 - The distribution component

Components of Docker

Docker Image

contains basic read-only image that forms the basis of container




Docker Registry

a repository of images which can be hosted publicly (like Docker Hub) or privately and behind a firewall



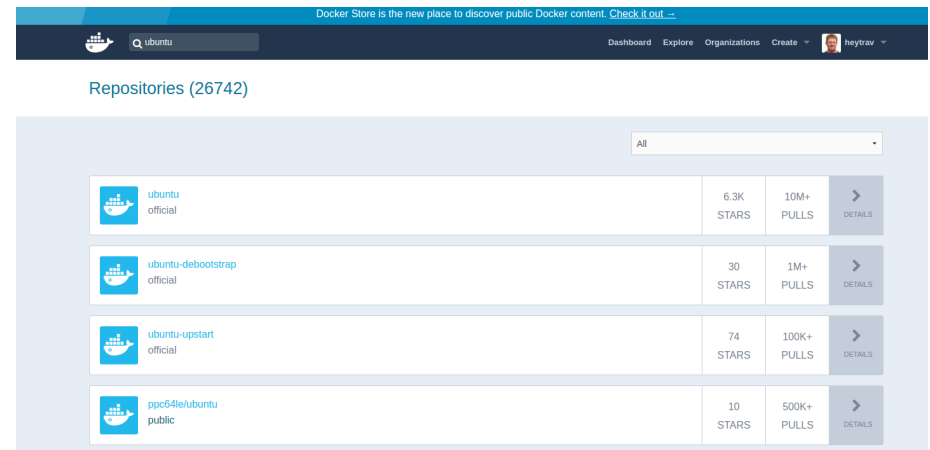
Docker Container

is comprised of a base image with layers that can be swapped out so it's not necessary to replace the entire VM when updating an application



Docker Registries

- Public repositories for docker images
 - **Docker Hub**
 - **Quay.io**
 - GitLab ships with docker registry
- Create your own private registry **docker/distribution**



The screenshot shows the Docker Store interface. At the top, there's a navigation bar with a search bar containing 'ubuntu' and links for Dashboard, Explore, Organizations, and Create. Below the navigation bar, the title 'Repositories (26742)' is displayed. A dropdown menu is set to 'All'. The main content area shows a table of repositories with the following data:

Repository	Stars	Pulls	Details
ubuntu/official	6.3K STARS	10M+ PULLS	> DETAILS
ubuntu-debootstrap/official	30 STARS	1M+ PULLS	> DETAILS
ubuntu-upstart/official	74 STARS	100K+ PULLS	> DETAILS
ppc64le/ubuntu/public	10 STARS	500K+ PULLS	> DETAILS

Underlying technology

Go

Implementation language developed by Google

Namespaces

Provide isolated workspace, or *container*

cgroups

limit application to specific set of resources

UnionFS

building blocks for containers

Container format

Combined namespaces, cgroups and UnionFS

Images and Containers

Docker images

- Images are the basis of containers
- An image is a *readonly* file system similar to tar archive
- *Distributable* artefact of Docker

Types of images

Official Base Image

Created by single authority (OS, packages):

- ubuntu:16.04
- centos:7.3.1611
- postgres

Base Image

Can be any image (official or otherwise) that is used to build a new image

Child Images

Build on base images and add functionality (this is the type you'll build)

Image naming semantics

- No upper-case letters
- Tag is optional. Implicitly *:latest* if not specified
 - postgres:*9.4*
 - ubuntu == ubuntu:*latest* == ubuntu:*16.04*
- If pushing to a registry, need url and username
 - If registry not specified, docker.io is default:
 - *docker.io/username/my-image* == *username/my-image*
 - *my.reg.com/my-image:1.2.3*
 - GitLab registry accept several variants:
 - gitlab.catalyst.net.nz:4567/<group>/<project>:tag
 - gitlab.catalyst.net.nz:4567/<group>/<project>/optional-image-name:tag
 - gitlab.catalyst.net.nz:4567/<group>/<project>/optional-name/optional-image-name:tag

Layering of images

- Images are *layered*
- Any child image built by adding layers on top of base
- Each successive layer is set of differences to preceding layer
- A layer is an instruction that
 - change filesystem
 - tells Docker what to do when run

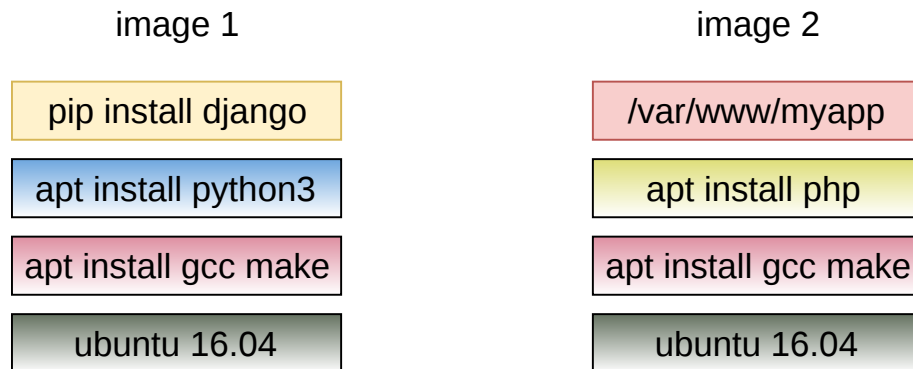
Layer	Description
4	execute <code>myfile.sh</code>
3	make <code>myfile.sh</code> executable
2	copy <code>myfile.sh</code> to working directory
1	install libs
0	Base Ubuntu OS

Sharing image layers

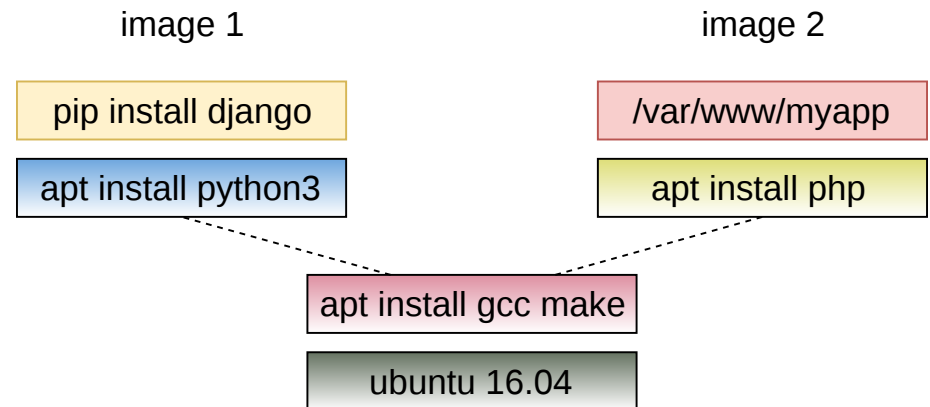
- Images will share any common layers
- Applies to
 - Images pulled from Docker
 - Images you build yourself

Sharing image layers

Two separate images



Reality: common layers shared



View image layers

`docker history <image>`


```
$ docker history heytrav/docker-introduction-slides
```

IMAGE	CREATED	CREATED BY	SIZE	COMMENT
e72084f25e08	2 months ago	/bin/sh -c #(nop)	0B	
<missing>	2 months ago	/bin/sh -c #(no	0B	
.				
.				
<missing>	9 months ago	/bin/sh -c #(n	0B	
<missing>	9 months ago	/bin/sh -c #(n	3.97MB	

Container basics


Namespaces

- Restrict visibility
- Processes inside a namespace should only see that namespace
- Namespaces:
 - pid
 - mnt
 - user
 - ipc




Cgroups

- Restrict usage
- Highly flexible; fine tuned
- Cgroups:
 - cpu
 - memory
 - devices
 - pids




Combining the two

A running container represents a combination of layered file system, namespace and sets of cgroups



Container layering

- Container creates its own read/write layer on top of image
- Multiple containers each have own read/write layer, but can share the actual image



Create images, explore layers

Docker command	Description	Syntax
diff	Inspect changes to files on a container's filesystem	<code>docker diff [options] <i>CONTAINERID</i></code>
commit	Create a new image from a container's changes	<code>docker commit [options] <i>CONTAINER</i> [<i>IMAGE[:TAG]</i>]</code>
history	Show history of an image	<code>docker history [options] <i>image:tag</i></code>

Exercise: Explore image layers

```
$ docker run -it ubuntu:16.04 /bin/bash
root@CONTAINERID:/$ apt-get update
root@CONTAINERID:/$ exit
$ docker ps -a
$ docker diff CONTAINERID
```

```
$ docker commit CONTAINERID ubuntu:update
13132d42da3cc40e8d8b4601a7e2f4dbf198e9d72e37e19ee1986c280ffcb97c
```

```
$ docker image ls
$ docker history ubuntu:16.04
$ docker history ubuntu:update
```

- Created an image by committing changes in a container
- Now have two separate images
- Share common layers; only difference is new layer on ubuntu:update

Creating Docker Images

Introducing the *Dockerfile*

- A text file
- Usually named `Dockerfile`
- Sequential instructions for building a Docker image
- Each instruction creates a layer on the previous

Structure of a Dockerfile

- Start by telling Docker which base image to use

```
FROM <base image>
```

- A number of commands telling docker how to build image

```
COPY . /app  
RUN make /app
```

- Optionally tell Docker what command to run when the container is started

```
CMD ["python", "/app/app.py"]
```

Common Dockerfile Instructions

FROM

FROM **image**:tag

Define the base image for a new image

```
FROM ubuntu:17.04
```

```
FROM debian # :latest implicit
```

```
FROM my-custom-image:1.2.3
```

- Image can be
 - An official base image
 - ubuntu:16.04
 - alpine
 - postgres:9.4
 - Another image you have created

RUN

RUN **command** **arg1** **arg2** ...

Execute shell commands for building image

```
RUN apt-get update && apt-get install python3
```

```
RUN mkdir -p /usr/local/myapp && cd /usr/local/myapp
```

```
RUN make all
```

```
RUN curl https://domain.com/somebig.tar | tar -xv | /bin/sh
```

COPY

COPY **src** **dest**

Copy files from build directory into image

```
COPY package.json /usr/local/myapp
```

```
COPY . /usr/share/www
```

WORKDIR

WORKDIR **path**

- Create a directory in the image
- Container will run relative to this directory

```
WORKDIR /usr/local/myapp
```

CMD

- Provide defaults to executable
- or provide executable
- Two ways to execute a command:
 - shell form: CMD **command** **param1** **param2** ...
 - exec form: CMD ["command", "param1", "param2"]

Exercise: Write a basic Dockerfile

```
$ cd ~/docker-introduction/sample-code/first-docker-file  
$ ls  
hello.sh
```

- Write a Dockerfile:
 - Named Dockerfile
 - Based on alpine
 - Set working directory to `/app`
 - Copy `hello.sh` into working directory
 - make `hello.sh` executable
 - tell docker to run `hello.sh` on docker run

```
FROM alpine  
WORKDIR /app  
COPY hello.sh .  
RUN chmod +x hello.sh  
CMD ["/hello.sh"]
```

Now that we have a Dockerfile, we can make an image

Building Docker images

`docker build [options] image:[tag] ./path/to/Dockerfile`

Options	Arguments	Description
<code>--compress</code>		Compress the build context using gzip
<code>-c, --cpu-shares</code>	int	CPU shares (relative weight)
<code>--cpuset-cpus</code>	string	CPUs in which to allow execution (0-3, 0,1)
<code>--cpuset-mems</code>	string	MEMs in which to allow execution (0-3, 0,1)
<code>--disable-content-trust</code>		Skip image verification (default true)
<code>-f, --file string</code>		Name of the Dockerfile (Default is 'PATH/Dockerfile')
<code>--pull</code>		Always attempt to pull a newer version of the image
<code>-t, --tag</code>	list	Name and optionally a tag in the 'name:tag' format

Note that a path to folder with Dockerfile is always required.
When in same directory, use "."

Exercise: build image using Dockerfile

- Build a Docker image:
 - Use Dockerfile from earlier example
 - Name image YOURNAME/my-first-image

```
$ docker build -t YOURNAME/my-first-image .
```

```
$ docker build -t YOURNAME/my-
```



00:06



```
$ docker run YOURNAME/my-first-image .
```

A few more Dockerfile directives

ENTRYPOINT

- Docker images need not be executable by default
- ENTRYPOINT configures executable behaviour of container
- *shell* and *exec* forms just like CMD

```
$ cd ~/docker-introduction/sample-code/entrypoint_cmd_examples  
$ docker build -t not-executable -f Dockerfile.notexecutable .  
$ docker run not-executable # does nothing
```

```
$ docker build -t executable -f Dockerfile.executable .  
$ docker run executable
```

Combining ENTRYPOINT & CMD

- Arguments following the image for `docker run image` overrides CMD
- Use exec form of ENTRYPOINT and CMD together to set base command and default arguments
- Hypothetical application

```
FROM ubuntu:latest
.  
.  
ENTRYPOINT ["/base-script"]  
CMD ["test"]
```

```
$ docker run my-image
```

By default this image will just pass `test` as argument to `base-script` to run unit tests by default

```
$ docker run my-image server
```

Passing argument at the end tells it to override CMD and execute with `server` to run server feature

Exploring ENTRYPOINT & CMD behaviour

- ```
$ cd sample-code/entrypoint_cmd_examples
```

- Compare Dockerfiles:
  - Dockerfile.cmd\_only
  - Dockerfile.cmd\_and\_entrypoint

- Build images:

```
$ docker build -t cmd_only -f Dockerfile.cmd_only .
$ docker build -t cmd_and_entrypoint -f Dockerfile.cmd_and_entrypoint .
```

- Run both the images with or without an additional argument to see what happens

# More Dockerfile instructions

## **EXPOSE**

ports to expose when running

## **VOLUME**

folders to expose when running

## **ENV**

Set an environment variable

See official reference [documentation](#) for more

# **Dockerising applications**

# Create web application in Docker

- Create a small web app based on Python Flask
- Write a Dockerfile
- Build an image
- Run the image
- Upload image to a Docker registry

# Step 1. Set up the web app

- Under `~/docker-introduction/sample-code/flask-app`

**app.py**

A simple flask application for displaying cat pictures

**requirements.txt**

list of dependencies for flask

**templates/index.html**

A jinja2 template

**Dockerfile**

Instructions for building a Docker image

# Our Dockerfile

```
FROM alpine:3.5

Install python and pip
RUN apk add --update py2-pip

install Python modules needed by the Python app
COPY requirements.txt /usr/src/app/
RUN pip install --no-cache-dir -r /usr/src/app/requirements.txt

copy files required for the app to run
COPY app.py /usr/src/app/
COPY templates/index.html /usr/src/app/templates/

tell the port number the container should expose
EXPOSE 5000

CMD ["python", "/usr/src/app/app.py"]
```



# Build the Docker image

```
$ cd ~/docker-introduction/sample-code/flask-app
$ docker build -t YOURNAME/myfirstapp .
```

```
→ docker build -t YOURNAME
```

Note: please replace YOURNAME with your Docker Hub username

# Run the container

```
$ docker run -p 8888:5000 --rm --name myfirstapp YOURNAME/myfirstapp
```

```
→ docker run -p 8888:5
```



00:06



...Now open **your test webapp**

# Login to a registry

```
$ docker login <registry url>
```

example-voting-app/vote

```
> docker login
```

Login with your Docker ID to push and pull images from Docker Hub. If you don't have a Docker ID, you can create one.

Username: heytrav

Password:

- If registry not specified, logs into [hub.docker.com](https://hub.docker.com)
- Can log in to multiple registries

# Push image to registry

```
$ docker push YOURNAME/myfirstapp
```

```
→ ~ docker push YOURNAME/my
```



00:06



# Summary

- Wrote a small web application
- Used Dockerfile to create an image
- Pushed image to upstream registry

# **Dockerfile best practices**

# General guidelines

- Containers should be as ephemeral as possible
- Avoid installing unnecessary packages
- Minimise concerns
  - Avoid multiple processes/apps in one container



# Use a `.dockerignore` file

```
.dockerignore
.git*
.dockerignore
Dockerfile
README*
don't import python virtualenv
.venv
```

- Top level of your project
- Very similar to `.gitignore`
- `COPY . dest/` will not copy files ignored in `.dockerignore`
- Include things you don't want in your image:
  - `.git` directory
  - `node_modules`, `virtualenv` directories

# General guidelines

- Use current official repositories in FROM as base image
- Image size may be a factor on cloud hosts where space is limited
  - debian 124 MB
  - ubuntu 117 MB
  - alpine 3.99 MB
  - busybox 1.11 MB
- Choice of image depends on other factors

# Layer caching

```
$ cd ~/docker-introduction/sample-code/caching
$ docker build -t caching-example -f Dockerfile.layering .
```

- Build image in sample-code/caching directory
- Run build a second time. What happens?
- Change line with Change me! and run again
- Each instruction creates a layer in an image
- Docker caches layers when building
- When a layer is changed Docker rebuilds from changed layer

# Consequences of layer caching

```
Example 1
FROM ubuntu:latest
RUN apt-get update
RUN apt-get install -y curl
#RUN apt-get install -y nginx
```

```
Example 2
FROM ubuntu:latest
RUN apt-get update \
 && apt-get install -y curl #nginx
```

```
$ cd ~/docker-introduction/sample-code/caching
$ docker build -t bad-apt-example -f Dockerfile.bad .
$ docker build -t good-apt-example -f Dockerfile.good .
```

- Uncomment nginx line and run `docker build` again
- Only rebuilds from layer that was *changed*
- Example 1: `apt-get update` does not refresh index
  - apt repos might change
- Best to combine `apt-get update` and install packages to force apt to refresh index (Example 2)

# Optimising image size

## Intermediate layers

- Image size is sum of intermediate layers
- Even if you remove something it exists as a diff on previous layer
- Run clean up in same layer whenever possible

# Example: Optimising image size

```
FROM ubuntu:latest 112MB
RUN apt-get update \
 && apt-get install -y \
 automake \
 build-essential \
 curl \
 wget \
 libcap-dev \
 reprepro 284MB
RUN rm -rf /var/lib/apt/lists/* 0MB
ADD https://dl.google.com/android/android-sdk_r24.4.1-linux.tgz . 326MB
RUN tar xf android-sdk_r24.4.1-linux.tgz 678MB
RUN rm -f android-sdk_r24.4.1-linux.tgz 0 MB
```

Image size: 1.4 GB

```
FROM ubuntu:latest 112MB
RUN apt-get update \
 && apt-get install -y \
 automake \
 build-essential \
 curl \
 wget \
 libcap-dev \
 reprepro \
 && rm -rf /var/lib/apt/lists/* 244MB
RUN wget https://dl.google.com/android/android-sdk_r24.4.1-linux.tgz && \
 tar xf android-sdk_r24.4.1-linux.tgz && \
 rm -f android-sdk_r24.4.1-linux.tgz 678 MB
```

Image size: 1 GB

# ADD

- Copies files to a directory

```
ADD . /usr/path/
```

- Downloads file from web

```
ADD http://domain.com/file.txt /usr/path/
```

- Unpack archives into directory

```
ADD file.tar /usr/path/
```

- However, does not unpack remote archives. This will just put `file.tar` in `/usr/path/`

```
ADD http://domain.com/file.tar /usr/path/
```

# Problems with ADD

- Large intermediate layers

```
ADD http://domain.com/big.tar.gz /usr/path/ # large intermediate layer
RUN cd /usr/path && tar -xvf big.tar.gz \
 && rm big.tar.gz
```

- Increased overall image size
- Better solution:

```
RUN curl -SL http://domain.com/big.tar.gz \
 | tar -xJC /usr/path
```

- Smaller image size
- COPY only copies files

```
COPY . /usr/path/
```

- Recommend to only use COPY and never ADD



# Multistage builds

## Optimise image builds

- Best practices intended to optimise image size by keeping them small
- Come at the expense of readability
  - Layers with long complicated commands
- Multistage builds
  - Introduced with Docker 17.05
  - Enable optimised image size
  - maintain readability

# Multistage builds

## How they work

- Multiple FROM directives in a Dockerfile
- Each FROM represents a new build
- Selectively copy artifacts from one of the previous builds
- Leave behind what is not needed

```
FROM ubuntu:16.04 as builder
WORKDIR /bin
COPY . /bin/
RUN make install

FROM alpine
COPY --from=builder /bin/myprogram /
ENTRYPOINT ['/root/myprogram']
```

## Multistage builds example

```
$ cd ~/href-counter
$ docker build -t href-counter -f Dockerfile.build .
$ docker image ls | grep href
```

| REPOSITORY   | TAG    | IMAGE ID     | SIZE  |
|--------------|--------|--------------|-------|
| href-counter | latest | b0eb64a75c55 | 687MB |

```
$ docker build -t href-counter-multi -f Dockerfile.multi .
$ docker image ls | grep href
```

| REPOSITORY         | TAG    | SIZE   |
|--------------------|--------|--------|
| href-counter-multi | latest | 10.3MB |

# CMD & ENTRYPOINT

## General best practices

- Avoid using *shell* form
  - `ENTRYPOINT "executable param1 param2 ..."`
- Docker directs POSIX commands at process with PID 1
- Using *shell* form, process is run internally using `/bin/sh -c` and do not have PID 1
- It can be difficult to stop container since process does not receive SIGTERM from `docker stop container`

```
$ cd ~/docker-introduction/sample-code/entrypoint_cmd_examples
$ docker build -t runtop-shell -f Dockerfile.top_shell .
$ docker run --rm --name topshell runtop-shell
```

What happens when you want to stop container *topshell*?

# CMD & ENTRYPOINT

## General best practices

- Best practice to use **exec** form:
  - `CMD ["executable", "param1", "param2", ...]`
- Or in form that creates interactive shell like
  - `ENTRYPOINT ["python"]`
  - `CMD ["/bin/bash"]`

```
$ docker build -t runtop-exec -f Dockerfile.top_exec .
$ docker run runtop-exec
```

- Sometimes app constraints don't allow single process on PID1
- For this purpose recommended to use **dumb-init**


# Summary

- Dockerfile best practices aim to
  - Keep image footprint small
    - Sometimes at expense readability
    - Multistage builds are a good compromise
  - Maintain clean control over containers
    - Easy to top and start
    - Flexible in the way they are executed
- Official Dockerfile **best practices**

# Docker and Development

# Microservices vs. Monoliths

- Small decoupled applications vs. one big app
- Developed independently
- Deployed and updated independently
- Scaled independently
- Better modularity
- Docker containers fit with microservice architecture





# Microservices in Docker

## Linking multiple containers

```
$ cd ~/docker-introduction/sample-code/mycomposeapp
```

- Let's build a simple application with two components
  - Web application using Python Flask
  - Redis message queue
- The app is already in mycomposeapp/app.py
- We want to run the app and redis as separate microservices
- Redis is already available as a **docker image**
  - `$ docker pull redis:alpine`
- We're going to have build a docker image for our app

# Create our app

- Go into the mycomposeapp directory

```
$ cd ~/docker-introduction/sample-code/mycomposeapp
$ ls
$ gedit Dockerfile
```

- Contents of Dockerfile

```
FROM python:3.4-alpine
WORKDIR /code
COPY requirements.txt /code
RUN pip install -r requirements.txt
COPY . /code
CMD ["python", "app.py"]
```

- Build Docker image for app

```
$ docker build -t web .
```

# Run containers as microservices

- First let's start our redis container

```
$ docker run -d --rm --name redis redis:alpine
```

- For our web container, we need a specific option to connect it to redis
  - `--link <name of container>`

```
$ docker run -d --rm --name web --link redis -p 5000:5000 web
```

- Once you start the web container go to **web page** to see counter

# Disadvantages of this approach

- Complicated with shell/script commands
  - Managing service interactions
  - Adding/managing services
- Can't scale services
- Stopping and cleaning up services can be tedious
  - BTW, you'll need to stop each of those containers

```
$ docker stop web
$ docker stop redis
```

- Better tools exist..

# Docker Compose

## Docker as a dev environment

- A tool that let's you easily bootstrap complex microservice apps
- Allows interactive development
  - you can work on the code while the container is running
- Can be used for staging/production environments
- Uses a YAML based config file called the *docker-compose file*

# The docker-compose file

- Service description file
- YAML
- By default: `docker-compose.yml`
- Specifies
  - Services
    - effectively containers that you will run
  - Volumes
    - filesystem mounts for containers
  - Networks
    - to be created and used by containers
- Have a look at the [compose file reference](#)

```

version: "3"
services:
 db:
 image: db
 volumes:
 - data-volume:/var/lib/db
 redis:
 image: redis:alpine
 webapp:
 build: .
 ports:
 - 80:80
 networks:
 hostnet: {}

volumes:
 data-volume:

networks:
 hostnet:
 external:
 name: host
```

# docker-compose services

- Each key in *services* dictionary represents a base name for a container
- Attributes of a service include
  - **build**
    - path or dictionary pointing to Dockerfile to build for container
  - **image**
    - Use a particular image for container
  - **ports**
    - expose ports for accessing application
  - **volumes**
    - mount into container

```

version: "3"
services:

 webapp:
 build: .
 ports:
 - 80:80
 - 443:443
 networks:
 hostnet: {}
 db:
 image: db
 volumes:
 - data-volume:/var/lib/db

 redis:
 image: redis:alpine
```

# Docker Compose

## Basic commands

`docker-compose` **COMMAND** [options] [args]

scale

| Command                | Description                          |
|------------------------|--------------------------------------|
| up                     | Start compose                        |
| down                   | Stop & tear down containers/networks |
| restart <service name> | Restart a service                    |

- Use `docker-compose -h` to view inline documentation



## Exercise: convert our app to use docker - compose

- In same directory as previous example
- Create a file called `docker-compose.yml`
- Add our service definition:

```

version: "3"
services:
 web:
 build: .
 ports:
 - "5000:5000"
 redis:
 image: redis:alpine
```

- Start our microservices

```
$ docker-compose up [-d]
```

# Scaling services

```
$ docker-compose up -d --scale SERVICE=<number>
```

- Try scaling the redis service to 4 instances

```
>> docker ps
```

| CONTAINER ID | IMAGE                              | COMMAND                  | CREATED |
|--------------|------------------------------------|--------------------------|---------|
| b11b11d8d02c | redis:alpine                       | "docker-entrypoint.s..." | 11 minu |
| 869f3b23c67c | mycomposeapp_web                   | "python app.py"          | 11 minu |
| ee70f4a2c2e2 | heytrav/docker-introduction-slides | "/usr/local/bin/dumb..." | 2 hours |

```
>> docker compos
```



00:06



# Stopping docker-compose

- In the directory where your docker-compose.yml file is:
- ```
$ docker-compose stop
```

Summary

- docker - compose provides useful way to setup development environments
- Takes care of
 - networking
 - linking containers
 - scaling services

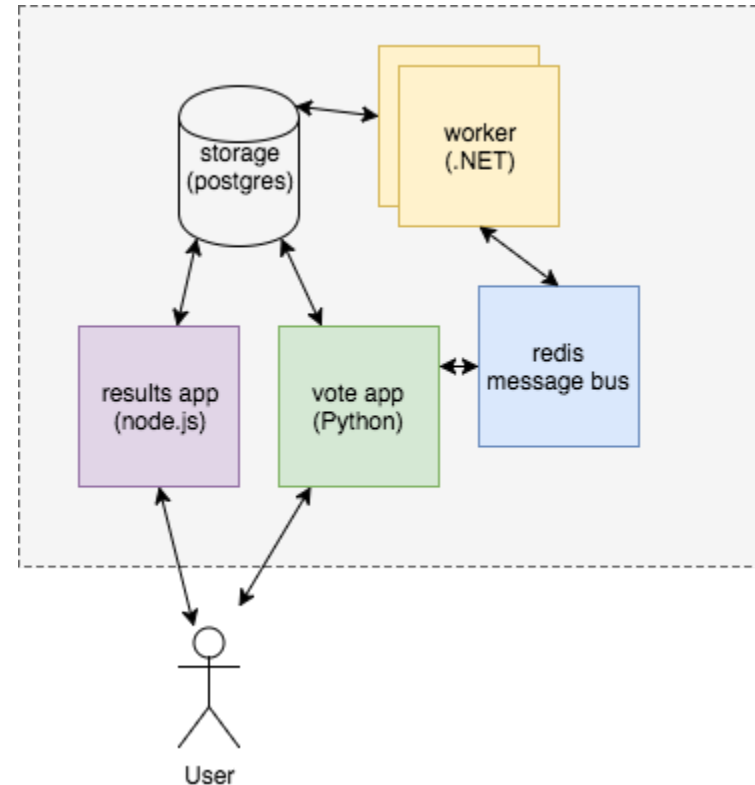
Docker Workflow Example

Taking apps from desktop to deployment

Example voting application

Microservice application consisting of 5 components

- Python web application
- Redis queue
- .NET worker
- Postgres DB with a data volume
- Node.js app to show votes in real time



```
$ git clone https://github.com/dockersamples/example-voting-app.git
$ cd example-voting-app
```

Start Application

```
$ cd example-voting-app  
$ docker-compose up -d
```

example-voting-app

```
> docker-compose up -d  
Creating network "examplevotingapp_front-tier" with the default driver  
Creating network "examplevotingapp_back-tier" with the default driver  
Building vote  
Step 1/7 : FROM python:2.7-alpine  
----> 9b06bbaac1c7  
Step 2/7 : WORKDIR /app  
----> 7d9f16f3c573  
Removing intermediate container acc64d316d13  
Step 3/7 : ADD requirements.txt /app/requirements.txt
```



00:06





Vote and view results

Interactive development

- Open up `vote/app.py`
- On lines 8 & 9, modify vote options
- View change in `voting` application

Change vote options

```
1 from flask import Flask, render_template, request, make_response, g
2 from redis import Redis
3 import os
4 import socket
5 import random
6 import json
7
8 option_a = os.getenv('OPTION_A', "Cats")
9 option_b = os.getenv('OPTION_B', "Dogs")
10 hostname = socket.gethostname()
11
12 app = Flask(__name__)
13
14 def get_redis():
15     if not hasattr(g, 'redis'):
16         g.redis = Redis(host="redis", db=0, socket_timeout=5)
17     return g.redis
18
19 @app.route("/", methods=['POST','GET'])
```

NORMAL   master  vote/app.py 

[Pymode] Activate virtualenv: /home/travis/workspace/catalystcloud-ansible/ansible-venv

Developer workflow

- Push code to repository
- Continuous Integration (CI) system runs tests
- If tests successful, automate image build & push to a docker registry
- Easy to setup with existing services
 - **DockerHub** (eg. **these slides**)
 - **GitHub**
 - **CircleCI**
 - **GitLab**
 - **Quay.io**

Developer workflow

- Ship your artefact directly using docker-compose
 - Useful if you want to test an image immediately
- Tell docker-compose to rebuild the image
- Let's build and tag the image as YOURNAME/vote:v2 and push to hub.docker.com
- This will come in handy in an example we're doing later

```
$ docker-compose build vote
$ docker tag examplevotingapp_vote:latest YOURNAME/vote:v2
$ docker push YOURNAME/vote:v2
```

Summary

- With docker-compose it's relatively easy to develop on a microservice application
- Changes visible in real time
- Can easily package and distribute images for others to use

Deploying Applications

Deploying Applications

- Inevitable goal of developing apps is to deploy them somewhere
- Typically some kind of hosting provider
 - Bare metal
 - Cloud Provider
- We'll use Catalyst Cloud OpenStack

First, some more buzzwords

- Immutable infrastructure
- Cattle vs pets
- Snowflake Servers **vs.** Phoenix Servers

Immutable Architecture/Infrastructure

- Phoenix servers
- The environment is defined in code
- If you need to change *anything* you create a new instance and destroy the old one
- Docker makes it much more likely you will work in this way



Orchestration

- Hosting may consist of multiple machines
- Once infrastructure is in place, need way to manage containers
 - networking
 - volume mounts
 - linking between containers
 - general monitoring, healthchecks
 - Deploying new images
- This is where orchestration tools come in

Container orchestration

- Frameworks for container orchestration
 - Docker Swarm
 - Kubernetes
- Manage deployment/restarting containers across clusters
- Networking between containers (microservices)
- Scaling microservices
- Fault tolerance

Kubernetes

- Container orchestrator
- Started by Google
- Inspired by Borg (Google's cluster management system)
- Open source project written in Go
- Cloud Native Computing Foundation
- Manage applications not machines


Docker Swarm

- Standard since Docker 1.12
- Manage containers across multiple machines
 - Scaling services
 - Healthchecks
 - Load balancing



Docker Swarm

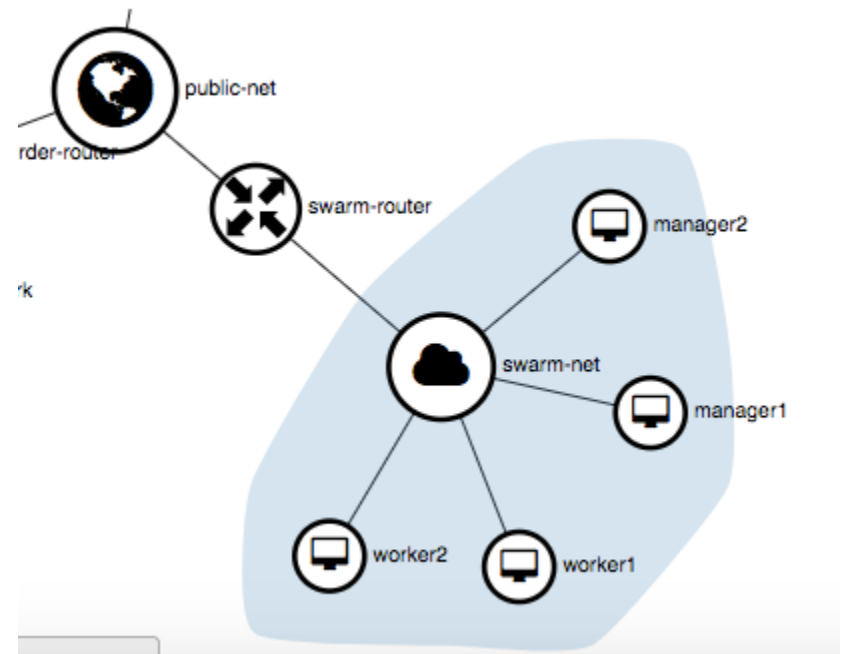
- Two types of machines or *nodes*
 - 1 or more *manager* nodes
 - 0 or more *worker* nodes
- Managers control global state of cluster
 - Raft Consensus Algorithm
 - If one manager fails, any other should take over



Deploying a Swarm Application

Setting up a cluster

- Need to:
 - provision machines
 - set up router(s)
 - set up security groups
- Preferable to use automation tools:
 - Chef
 - Puppet
 - Terraform
 - Ansible



Create a cluster

```
$ cat ~/credentials.txt
$ source ~/os-training.catalyst.net.nz-openrc.sh
<enter os training password from ~/credentials.txt>

$ cd ~/docker-introduction/ansible
$ ansible-playbook -i cloud-hosts -K -e suffix=-$(hostname) \
  create-swarm-hosts.yml
```

This will do stuff for while, good time for some coffee

Create Swarm

```
$ ssh manager<TAB><ENTER>  
$ docker swarm init
```

```
ubuntu@manager1-trainingpc:~$ docker swarm init
```



00:06



Copy the `docker swarm join ...` command that is
output

Join Worker Nodes

Paste the command from the manager node onto command line.

```
$ ssh worker1<TAB><ENTER>  
$ docker swarm join --token $TOKEN 192.168.99.100:2377
```

```
ubuntu@worker1-trainingpc:~$ docker swarm join \  
> --token SWMTKN-1-12ffzfi8ecbk7420j3ill7u0fwww3lqrm63egsyznftv0rp99r-7j6k8emsis34yl2  
> 192.168.99.100:2377  
This node joined a swarm as a worker.  
ubuntu@worker1-trainingpc:~$
```



00:03

Repeat this for worker2

Check nodes

```
$ docker node ls
```

```
ubuntu@manager1-trainingpc:~$
```



00:00



Swarm Stack File

- Service description
- YAML format
- Similar to file used for docker-compose
- A few differences
 - No build option
 - No shared volumes

```
# stack.yml
version: "3.3"
services:
  db:
    image: postgres:9.4
    .
    .
  redis:
    image: redis:latest
    deploy:
      replicas: 3
  vote:
    image: vote:latest
    depends_on:
      - redis
      - db
    denlov:
```

Deploying voting app

Upload docker-stack.yaml to manager node

```
$ cd ~/example-voting-app  
$ scp docker-stack.yaml manager-TRAININGPC:~/
```

Deploy application

```
$ docker stack deploy -c docker-stack.yml vote
```

Monitor deploy progress

```
$ watch docker stack ps vote
```

```
$ watch docker service ls
```


Try out the voting app

<http://voting.app:5000>

To vote

<http://voting.app:5001>

To see results

<http://voting.app:8080>

To visualise running containers

Scale services

```
$ docker service scale vote_vote=3
```

Look at the changes in the **visualizer**

Update a service

```
$ docker service update --image YOURNAME/vote:v2 vote_vote
```

```
ubuntu@manager1-trainingpc:~$ docker service update --image
```



00:06



Now go to the **voting app** and verify the change

Drain a node

```
$ docker node update --availability drain worker1
```

- Sometimes necessary to take host offline
 - Planned maintenance
 - Patching vulnerabilities
 - Resizing host
- Prevents node from receiving new tasks
- Manager stops tasks running on node and launches replicas on active nodes

Return node to service

```
$ docker node update --availability active worker1
```

- during a service update to scale up
- during a rolling update
- when you set another node to Drain availability
- when a task fails on another active node

Tear down your cluster

When you're done playing around with the voting app, please run the following

```
$ ansible-playbook -i cloud-hosts -K -e suffix=-$(hostname) \
  remove-swarm-hosts.yaml
```

Summary


- Created a cluster with a cloud provider using ansible
 - 1 manager node
 - 2 worker nodes
- Deployed microservice for voting app in Docker Swarm
- Scaled service from 2 to 3 services
- Rolling-Updated image

Wrap up

Docker ecosystem

- An explosion of tools
- Hard to keep up
- Lets have a quick look

Docker ecosystem



Competing technologies

- rkt (CoreOS)
- Serverless (FaaS)
 - Lambda (AWS)
 - Azure Functions (Microsoft)
 - Google cloud functions
 - iron.io

The end